

Model evaluation

Marko Medved

I. CROSS-VALIDATION

Cross-validation was implemented and evaluated on four different models. The performance of each model was assessed using log loss and accuracy as evaluation metrics.

A. Implementation details

Since the dataset is assumed to be representative of the data generating process and the target class distribution is imbalanced, stratified cross-validation was implemented. This ensures that each fold maintains the same class distribution as the overall dataset, preventing the learning algorithm from encountering folds where certain classes are underrepresented. Preserving class balance ensures the model can learn patterns from all classes, leading to more reliable evaluation.

As instructed, a baseline classifier and logistic regression were first evaluated. For the third model, a decision tree was chosen due to its sensitivity to the minimum number of observations required to split a node — a key parameter that was optimized within each fold.

The parameter optimization was carried out in two ways:

- **Training Fold Optimization:** Trained on the training split of each fold; the parameter that produced the lowest log loss was selected.
- **Nested Cross-Validation:** Applied an additional layer of cross-validation within the training split. The parameter with the lowest cumulative log loss from the inner folds was selected.

Log loss was chosen as the evaluation metric instead of accuracy because it is a strictly proper scoring rule — meaning it encourages the model to output the true class probability distribution as opposed to accuracy where only the mode of the distribution is taken into account.

The choice of $k = 10$ for cross-validation was made because decision trees are inherently unstable models — they are sensitive to small variations in the training data, which can increase variance when using a larger k . Since the dataset is reasonably large, 10 folds strike a good balance between bias and variance, providing a stable estimate of model performance without introducing excessive bias from smaller training sets.

A similar reasoning applies to the inner loop of the nested cross-validation. With only 10% less data available in the inner loop compared to the outer loop, using the same value of k (10) should still maintain an effective balance between training size and validation stability.

B. Results

Table I shows the performance of the four models. Note that the uncertainty was quantified using bootstrap with 500 repetitions.

Focusing on log loss, the decision tree shows a particularly high value, especially when using training fold optimization. This can be attributed to the fact that this model is often overly confident in its predictions, meaning the predicted class tends to have a very high probability. Consequently, in the case of misclassification, this high confidence leads to a larger penalty, resulting in high log loss. This effect is especially

evident when using training fold optimization, where the tree overfits more severely than in the nested CV case. Note that the smallest possible minimal number of observations required to split a node during parameter tuning was set to 10, which helps prevent the tree from completely overfitting.

When examining accuracy, the results are more consistent when compared to logistic regression, which is expected since accuracy only considers the predicted class and not the full predicted distribution.

Model	Log loss	Accuracy
Baseline	1.161 ± 0.013	0.6121 ± 0.0069
Logistic Regression	0.672 ± 0.013	0.7329 ± 0.0060
Decision Tree	2.569 ± 0.095	0.7217 ± 0.0063
Decision Tree (Nested CV)	1.378 ± 0.059	0.7386 ± 0.0060

Table I

PERFORMANCE COMPARISON OF MODELS

II. EFFECT OF DISTANCE ON ERROR

To investigate the influence of distance on errors, two different methods were applied.

Firstly, a visual method involved plotting a scatter plot of log losses at different distances, along with a fitted linear regression trend shown in Figure 1. The plot includes results for all models, where we observe a general trend of decreasing log loss at larger distances. For the baseline model, the plot may appear unusual; however, this is expected since the number of distinct log losses is limited by the number of shot types.

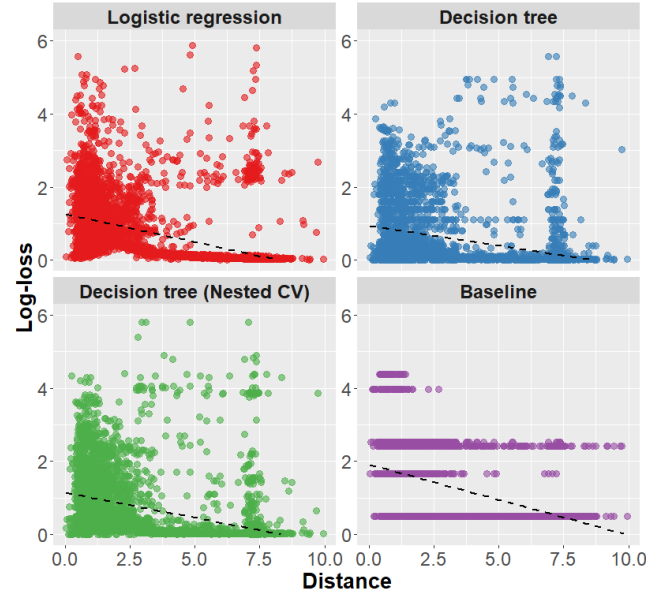


Figure 1. Effect of Distance on Log-Loss Across Models

The second method involved calculating the correlation coefficient between distance and log loss to quantify their linear relationship. As shown in Table II, all four models produced a negative correlation coefficient, indicating that log loss tends to decrease with increasing distance, consistent with the visual observation.

Model	Correlation
Baseline	-0.562 ± 0.009
Logistic Regression	-0.455 ± 0.018
Decision Tree	-0.311 ± 0.009
Decision Tree (Nested CV)	-0.245 ± 0.008

Table II

CORRELATION BETWEEN DISTANCE AND LOG-LOSS ACROSS MODELS

This relationship can be explained by examining the distribution of shot types. Figure 2 shows that at higher distances, most shots are classified as above head shots, whereas the distribution is more balanced in the overall dataset and at lower distances. For high-distance shots, the top 25% of distances were selected, while for low-distance shots, the bottom 25% were used. This also explains why the baseline model exhibits the highest correlation, since it always assigns the highest probability to the majority class.

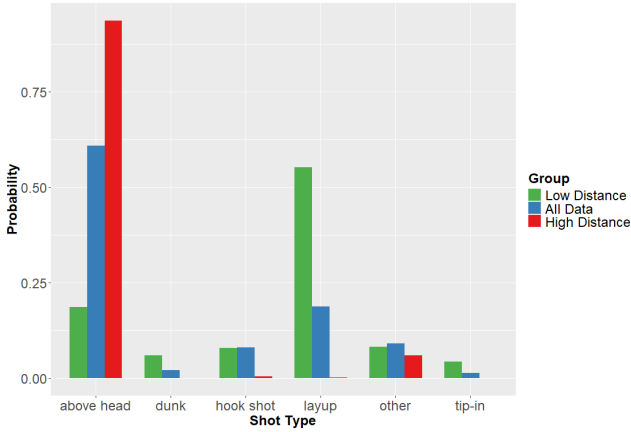


Figure 2. Shot Type Distribution By Distance

III. ESTIMATING PERFORMANCE ON TRUE COMPETITION DISTRIBUTION

A. Implementation details

To estimate how our model would perform on the true distribution of competition types, a weighted bootstrap method was implemented.

Unlike standard bootstrap, where samples are drawn uniformly, the weighted bootstrap adjusts the sampling probabilities based on the difference between the true and observed class distributions. Specifically, the sampling weights were calculated by taking the ratio of the true class probability (for competition types) to the current class probability. The weights were then normalized by dividing by their sum to create a valid probability distribution.

This approach ensures that the bootstrap samples reflect the true distribution of competition types, providing a more accurate estimate of model performance. Note that this method was applied directly to the log-losses and accuracies calculated in the first part.

B. Results

Table III shows the results of this estimation. The results differ slightly from the original ones, with most models showing a slight improvement — particularly the baseline model, which demonstrates noticeably better performance.

Model	Log loss	Accuracy
Baseline	1.126 ± 0.013	0.6351 ± 0.0068
Logistic Regression	0.652 ± 0.013	0.7525 ± 0.0061
Decision Tree	2.599 ± 0.094	0.7298 ± 0.0063
Decision Tree (Nested CV)	1.368 ± 0.060	0.7490 ± 0.0062

Table III

PERFORMANCE COMPARISON ON TRUE COMPETITION DISTRIBUTION

To explain this, we need to revisit the distribution of shot types — this time in relation to the competition type, as shown in Figure 3. By adjusting the data to reflect the true distribution, we increased the influence of data points from the NBA competition type while reducing the influence of others.

Since the above head shot type is more frequent in NBA data points, it is not surprising that the performance of the baseline model improved.




Figure 3. Shot Type Distribution by Competition